

Qualifying the Measurement Instrument

Phase 0 of the Phase 1 replacement programme — and five occasions on which a measurement artefact was read as a result about the method

August 2026

Abstract

Before approaches **B** (auction-dynamics MBO) and **C** (capacity-constrained geodesic Lloyd) can be compared against the PGD control on perimeter, the things that would do the comparing must themselves be shown to work: a *balanced-assignment solver* shared by all three approaches, and a single *evaluation harness*. This report records that qualification. Four results. (1) The harness reproduces the validated $N=100$ deliverable's Phase 2 perimeter **bitwise** — 185.2546144718457 — from a completely different entry path, which simultaneously establishes that hardening an arm's labels costs it nothing and that Phase 2 is bit-deterministic on this host. (2) The shared solver needed *no new algorithm*: the conclusion that an auction or Sinkhorn solver was required came from a fixture modelling a state the arms occupy for one outer iteration. (3) Under the final protocol (three meshes $\times$ two arms $\times$ three seeds) the solver passes **18/18**, and the gate is shown to discriminate — the pre-0a un-normalized solver is *indistinguishable from doing nothing*. (4) Measured at genuinely equal footing the arms sit at **parity**, $\pm 10\%$, with the incumbent, not clear of it; two earlier versions of this claim were wrong in the arms' favour. The methodological content is larger than the technical content: **five** times during this work a defect in the measurement was read as a result about the method, and the report documents what caught each one.

Provenance

Status	measured
Date	2026-08-16 to 2026-08-18
Surface	torus $T(R=1.0, r=0.6)$ throughout
Fixtures	$V=47,488$ (224×212) at $N=300$; $V=114,144$ (348×328) at $N=100$ and at $N=300$ (the production configuration). Seeds 0, 1, 2.
Control	<code>results/run_20260709_081548_surftorus_npart100_..._lam5.1_seed84172851</code> , the validated $N=100$ deliverable: Phase 2 best perimeter 185.2546144718457 at iteration 20 of 20, campaign <code>ipopt.btol0.001_lbfgs30_hess_bestiter_partial</code> , worst cell 0.7786%, 0 dead / 0 weak / 0 fragmented.
$N=300$ baselines	<code>run_20260806_123326</code> ($V=47,488$, Phase 2 best 323.319247087254, plateaued at iterate 18 of 19) and <code>run_20260808_191030</code> ($V=114,144$, 322.96218780465847, <i>censored</i> — best is its final iterate). These are <i>one</i> experiment: the second resumed from the first (<code>warm_start_path</code>). Both were copied into this checkout on 2026-08-18 (12.6 GB, verified identical).
Incumbent	approach A's own dual stage on its own log densities, normalized, at the same 2000-iteration budget the arms are scored with: 1.3517% at $V=47,488$, 0.5991% at $V=114,144$, both $N=300$.
Scripts	<code>testing/test_balanced_assignment_solver.py</code> (gates 1–4); <code>testing/test_arm_harness.py</code> (0b acceptance); <code>testing/test_phase0d_shakedown.py</code> (0d, ~ 34 min); <code>docs/experiments/07-phase0-shared-harness/make_figures.py</code> .
Code	<code>src/partition/balanced_readout.py</code> (<code>solve_dual_offsets</code> , <code>assignment_margin_scale</code>); <code>src/partition/arm_harness.py</code> (<code>assignment_quality_bar</code> , <code>A_PARITY_NORMALIZED</code>).
Versions	Python 3.12.13, numpy 2.4.4, scipy 1.17.1, h5py 3.16.0; macOS 15.3.1 arm64 (Mac mini).
Anchors	0d reproduces 185.2546144718457 bitwise ; final gate 2 is 18/18 in 2 h 39 m; the null solver scores 16.588% at the production pin against a 0.841% bar.

1 Where B and C come from

Neither arm is invented here; both are established constructions applied to this surface and this pipeline. This section states the provenance explicitly, because an earlier draft of this report carried none.

B — volume-constrained threshold dynamics. Threshold dynamics originates with Merriman, Bence and Osher [14]: alternate short-time diffusion of the phase indicators with a pointwise reassignment, and the interfaces move by (approximate) mean curvature. Esedoğlu and Otto [7] put the multiphase case on a variational footing, exhibiting a Lyapunov functional that Γ -converges to a weighted perimeter — the same limit our own energy targets [3, 15]. The specific variant we would build is Jacobs, Merkurjev and Esedoğlu's *auction dynamics* [10], which imposes exact volume constraints by replacing the pointwise threshold with a *balanced assignment*, solved by Bertsekas' auction algorithm [2]. That paper is the direct antecedent: it is volume-constrained multiphase MBO, published for this problem class.

Two caveats the literature itself supplies. Connectivity is *not* guaranteed per step — stray components shrink under curvature flow but are not excluded — and on graphs and coarse meshes threshold dynamics is known to *freeze* when the diffusion length falls to the discretization scale [16]. Both shaped the design: the first sets the validity grading, the second fixes $\tau = (4h)^2$ rather than h^2 .

C — capacity-constrained centroidal power diagrams. Lloyd’s algorithm [12] iterates “assign to nearest site, move site to centroid”. Adding a prescribed-capacity constraint turns the assignment into a semi-discrete optimal transport problem whose solution is a *power diagram* with per-site weights — the Minkowski-type result of Aurenhammer, Hoffmann and Aronov [1]. Mérigot [13] and Kitagawa, Mérigot and Thibert [11] give the modern solvers for those weights; de Goes *et al.* [6] use exactly this construction for blue-noise sampling. On a curved surface the Euclidean distance is replaced by a geodesic one, for which the heat method [4] is the standard alternative to the graph-Dijkstra approximation used in our fixtures.

C minimizes a *quantization* energy, not perimeter. It lands near the perimeter optimum only because both optima are asymptotically hexagonal in the plane — Hales for perimeter [9], Fejes Tóth and Gershgorin [8] for quantization. **Both are flat-plane results**, and this torus has Gaussian curvature of both signs, so neither transfers; how far C lands from optimal is precisely what remains to be measured.

The shared solver, and what was nearly rebuilt. All three approaches need the same primitive, and it is the semi-discrete OT dual [1, 13]: find per-cell offsets ψ such that $\arg \max_k [s_{ik} + \psi_k]$ has equal masses. The shipped implementation is subgradient ascent on that dual. Section 4 records that this report’s programme nearly replaced it with an auction [2] or an entropic (Sinkhorn) solver [5] — both standard, either defensible — on evidence that turned out to be a fixture artefact. Sinkhorn was implemented and measured; it *lost* to the incumbent on five of six score matrices.

2 Why Phase 0 exists

The Phase 1 relaxation costs **22.0–79.3 h** for a completed $N=300$ solution and its winner-take-all readout is not a valid partition, so it is repaired afterwards. That does not scale. Approaches B and C would produce a valid partition at every step in minutes. Deciding between them requires measuring *perimeter after Phase 2*, and both — plus the shipped approach A — need one shared primitive: *balanced assignment given arbitrary per-vertex scores*, which exists inside the balanced readout as `solve_dual_offsets`.

Phase 0 makes that primitive usable on foreign score scales (0a), builds one evaluation harness (0b), adopts the $N=300$ baselines (0c), and proves the harness reproduces a known answer (0d). **It measures nothing about perimeter**; every question about whether B or C is any good remains open.

3 0d — the harness reproduces a known answer, bitwise

The strongest single result. The $N=100$ control’s own winner-take-all labels were hardened to one-hot densities, written in the Phase 1 solution schema, and pushed through the entire harness under the control’s own campaign `refinement.yaml`:

reference	185.2546144718457	
harness	185.2546144718457	(iterate 20 of 20)
relative difference	0.000000%	bitwise: true wall 2023 s

Independently verified: the scratch campaign holds 20 fresh iteration files, its `refinement.yaml` names the one-hot arm file as `base_solution`, and the *entire* 20-iterate trajectory is bitwise equal to the control campaign computed five weeks earlier, including the 14,883-entry `lambda_parameters` array.

Three consequences. **(a)** The apples-to-apples question is settled exactly: an arm’s hard labels

and PGD’s continuous field produce identical Phase 2 initial states, because `ContourAnalyzer` reads densities only through `argmax` and every variable point starts at $\lambda=0.5$. **(b)** Phase 2 is bit-deterministic on this host, previously an unverified load-bearing assumption. **(c)** The control is itself *censored* — its best iterate is its last — so 185.2546144718457 is a *bound*.

4 0a — no new solver was needed

`solve_dual_offsets` is subgradient ascent on the semi-discrete OT dual. It is *correct* for any additive score matrix but its step schedule was calibrated to the log-density scale. Normalizing scores by a robust assignment-margin scale, and raising the step (`normalized_eta0` 0.5 $\rightarrow$ 2.0; measured 3.654 $\rightarrow$ 2.209 $\rightarrow$ **1.672** $\rightarrow$ 2.248% for $\eta_0 = 0.5, 1, 2, 4$) and the budget (400 $\rightarrow$ 2000) is the whole fix. Approach A’s path is untouched and byte-identical.

4.1 The conclusion that was withdrawn

The first version of gate 2 scored a *cold start* — balance a raw geometric guess in one shot — and B failed at 45.35%. This was written up as structural: the shipped solver is subgradient ascent tuned for a small correction, B and C need to rebalance ~ 230 of 300 cells, therefore build a Bertsekas auction or an entropic solver. **All of that was wrong.** B and C are *iterative*: from outer iteration 1 each assignment starts from an already-balanced partition. An iterated Lloyd loop with the *existing* solver runs 23 consecutive iterations under the bar. Sinkhorn, implemented and measured, *loses* to the incumbent on 5 of 6 score matrices. Two knobs dismissed as dead — budget and step size — each fix it on their own; the “a $10\times$ budget costs 20 minutes” figure was operations *counted*, never timed (the true cost is 33.9 ms/iteration).

5 The gate, and whether it discriminates




Figure 1: Final protocol: three fixtures, two arms, three seeds, 18/18 pass. Red line: the acceptance bar. Grey dashed: the incumbent at equal budget.

The bar is $\max(2 \times \text{vertex granularity}, \text{A's own dual stall at equal budget})$. A bar the shipped solver itself fails cannot be evidence about a foreign score scale; a bar below the one-vertex granularity floor is impossible by construction.




Figure 2: Degraded solvers at the production pin, C scores. The pre-0a un-normalized solver scores 16.588% — *equal to the null solver to three decimals*, because its best iterate is its $\psi=0$ starting point.

The shipped configuration sits 27–48% under its bars with the failure boundary a factor 2–5 away. Gate 3 asserts a null solver *fails*; gate 4 verifies that early stopping is safe inside an iterated loop (3.0× faster for +0.008 points of degradation).

6 The result, stated at equal footing

fixture	bar	C (s0/s1/s2)	B (s0/s1/s2)
V=47,488, N=300	2.022%	1.472 / 1.392 / 1.495	1.405 / 1.390 / 1.406
V=114,144, N=100	0.280%	0.146 / 0.142 / 0.144	0.161 / 0.151 / 0.138
V=114,144, N=300	0.841%	0.624 / 0.573 / 0.609	0.537 / 0.573 / 0.607

Against the incumbent’s 0.5991% at the production pin, C spans 0.573–0.624% and B spans 0.537–0.607%: **parity**, $\pm 10\%$. Earlier drafts of this claim said “2–3× better” and then “clear of parity by 30%”; both compared arms at 2000 iterations against A at 400. *A budget difference wearing a method difference’s clothes*. The assignment step is not a differentiator and was never meant to be one — 0a’s job was to make the shared primitive usable so that perimeter can be measured.

7 Five measurement artefacts read as method results

1. **The stall guard** warned on any floored line search — which is how a healthy level normally ends. Two commits reasoned from the false positive.
2. **A corpus glob** reached 130 of 350 traces, and the resulting “empty gap” became a report’s strongest argument.
3. **A baseline audit** enumerated one working directory and concluded no $N=300$ Phase 2 baseline existed. Two did, in a sibling worktree.

4. **The theatre bar:** gate 2’s bar included $1.25\times$ a strong reference computed *with the solver under test*, so a broken solver inflated its own bar. Six crippled solvers passed, the weakest doing nothing at all.
5. **The early-stop mask:** with early stopping inside the measurement path, gate 2 passed 12/12 while every number hugged its bar from below (Figure 3).



Figure 3: Artefact 5. A passing gate looks identical either way; only the pre-registered prediction distinguished them.

Two things caught these, and neither was insight. **Pre-registration:** the early-stop mask surfaced only because predicted values were committed before the run, and four “missed” in a suspiciously uniform direction. **Asking whether anything fails:** every test asked “*does the shipped solver pass?*” and none asked “*does anything fail?*”, so a gate that passed everything looked exactly like a gate that worked. Gate 3 now closes that.

On “survived adversarial review”. Four adversarial reviews inform this work and each found real defects. They were subagents dispatched *from the same session as the work they reviewed* — independent of the author’s reasoning, not of the author’s framing, tooling, or choice of what to point them at. Every load-bearing number they produced has since been self-reproduced, except one marked indicative in the plan. That caveat belongs next to the claim.

8 Conclusions

1. The harness is qualified: it returns a known answer bitwise, and hardening an arm’s labels costs nothing measurable.
2. No new assignment solver is needed. Normalization, step size and budget suffice; the auction/Sinkhorn conclusion was a fixture artefact.
3. The gate discriminates: 18/18 for the shipped configuration, with the pre-0a solver indistinguishable from doing nothing.
4. Arms and incumbent are at parity $\pm 10\%$ on assignment quality. The open question — whether B or C produces *shorter perimeter* — is untouched by this report.

Related documents

- Plan and pre-registration: docs/plans/PHASE1_BC_REPLACEMENT_PLAN.md
- Taxonomy: docs/reference/PHASE1_HIGHN_APPROACHES_ABCDE.md
- The gap being closed: docs/reference/winner_take_all_partition_gap.md
- Prior arms: docs/experiments/05-soft-area-constraint/, docs/experiments/06-subfloor-ladder/

References

- [1] Franz Aurenhammer, Friedrich Hoffmann, and Boris Aronov. Minkowski-type theorems and least-squares clustering. *Algorithmica*, 20(1):61–76, 1998. Power diagrams with weighted sites realise constrained least-squares assignments with prescribed cluster sizes: the capacity-constrained assignment underlying the balanced readout’s offsets and approach B’s initialisation. Held since 2026-09-23; authors, title, journal, volume 20 (1998), pages 61–76 verified from the first page.
- [2] Dimitri P. Bertsekas. The auction algorithm: A distributed relaxation method for the assignment problem. *Annals of Operations Research*, 14(1):105–123, 1988.
- [3] Benjamin Bogosel and Édouard Oudet. Partitions of minimal length on manifolds. *Experimental Mathematics*, 26(4):496–508, 2017. THE foundational reference for this project: Phase 1 (Γ -convergence relaxation) and Phase 2 (direct perimeter minimization). Distinct from bogosel2021longest, which is the same authors on the OUTER problem and is cited only in the literature section of experiment 08 – do not conflate them. read year = 2023 with no volume or pages, under the key bogosel2023partitions. The rest of the repository already said 2017 (src/surfaces/double_torus.py, banchoff_chmutov.py, docs/plans/MESH_DEGENERACY_AND_NEEDLE_TRIANGLES.md); the bibliography was the outlier. Taylor & Francis lists this as *Experimental Mathematics volume 26 (2017), issue 4*, and the article page gives “Published online: 04 Oct 2016”. Crossref carries both dates and its automatic BibTeX transform emits year = 2016; the DOI also encodes 2016. We cite the ISSUE year, 2017. A venue preferring online-first dating would write 2016; both refer to this same article, DOI 10.1080/10586458.2016.1223570, arXiv:1606.02873.
- [4] Keenan Crane, Clarisse Weischedel, and Max Wardetzky. Geodesics in heat: A new approach to computing distance based on heat flow. *ACM Transactions on Graphics*, 32(5):152:1–152:11, 2013.
- [5] Marco Cuturi. Sinkhorn distances: Lightspeed computation of optimal transport. In *Advances in Neural Information Processing Systems*, volume 26, pages 2292–2300, 2013. Entropic regularisation of optimal transport. Implemented and measured for this project’s balanced assignment, where it lost to subgradient ascent on five of six score matrices (docs/experiments/07-phase0-shared-harness/). Held as docs/papers/cuturi_2013.sinkhorn.distances.pdf since 2026-09-23 (the NeurIPS proceedings PDF); author and title verified from it. Volume and page range are from the proceedings listing — the paper PDF carries no page numbers.
- [6] Fernando de Goes, Katherine Breeden, Victor Ostromoukhov, and Mathieu Desbrun. Blue noise through optimal transport. *ACM Transactions on Graphics*, 31(6):171:1–171:11, 2012.
- [7] Selim Esedoğlu and Felix Otto. Threshold dynamics for networks with arbitrary surface tensions. *Communications on Pure and Applied Mathematics*, 68(5):808–864, 2015.

- [8] Allen Gersho. Asymptotically optimal block quantization. *IEEE Transactions on Information Theory*, 25(4):373–380, 1979.
- [9] Thomas C. Hales. The honeycomb conjecture. *Discrete & Computational Geometry*, 25(1):1–22, 2001.
- [10] Matt Jacobs, Ekaterina Merkurjev, and Selim Esedoğlu. Auction dynamics: A volume constrained MBO scheme. *Journal of Computational Physics*, 354:288–310, 2018. AUTHORS CORRECTED 2026-08-21, verified against dblp (journals/jcphy/JacobsME18) and ScienceDirect. This entry previously read “Jacobs, Kim and Léger”, an error that originated in the archived A–E proposal (docs/reference/phase1_highn_proposal_ABCDE.original.md) and propagated into seven documents before being checked. Volume, pages and year were correct throughout. Note Esedoğlu is an author of both this and `esedoglu2015threshold`.
- [11] Jun Kitagawa, Quentin Mérigot, and Boris Thibert. Convergence of a Newton algorithm for semi-discrete optimal transport. *Journal of the European Mathematical Society*, 21(9):2603–2651, 2019. Damped Newton for the semi-discrete dual; exploits smoothness in the continuous-source setting, which is why it does not transfer to this project’s finite-to-finite problem (see docs/math/09-balanced-readout). Held since 2026-09-23 — the HAL preprint (hal-01290496), not the published version; all fields verified from its citation block.
- [12] Stuart P. Lloyd. Least squares quantization in PCM. *IEEE Transactions on Information Theory*, 28(2):129–137, 1982.
- [13] Quentin Mérigot. A multiscale approach to optimal transport. *Computer Graphics Forum*, 30(5):1583–1592, 2011. Multiscale solver for semi-discrete optimal transport. Held since 2026-09-23 — the HAL preprint (hal-00604684), not the published version; author, title, journal, volume 30(5), year and DOI verified from its citation block. PAGE RANGE DISCREPANCY: the HAL block says 1584–1592, this entry and Crossref say 1583–1592. Not resolvable from the preprint; check the published article before external distribution.
- [14] Barry Merriman, James K. Bence, and Stanley J. Osher. Motion of multiple junctions: A level set approach. *Journal of Computational Physics*, 112(2):334–363, 1994.
- [15] Luciano Modica and Stefano Mortola. Un esempio di γ -convergenza. *Bollettino dell’Unione Matematica Italiana B (5)*, 14:285–299, 1977. The original Gamma-convergence example behind the Phase 1 relaxation energy. NOT HELD — searched 2026-09-23 without success; in Italian, and the Bollettino UMI of that period is poorly digitised. Cited here for priority, not for content: the result this project actually uses (the interface constant $c_W = 2 \int_0^1 \sqrt{W}$) is proved in English in `modica1987gradient`, which IS held, and the nonlocal analogue in `alberti1998nonlocal`. Every field above is from the standard citation, not from a copy. Do not spend more time looking for it.
- [16] Yves van Gennip, Nestor Guillen, Braxton Osting, and Andrea L. Bertozzi. Mean curvature, threshold dynamics, and phase field theory on finite graphs. *Milan Journal of Mathematics*, 82(1):3–65, 2014.